


Heapsort



Consider the sorting problem discussed in [another chapter](#). Namely, consider the problem of [permuting](#) the elements of an array $v[1 \dots n]$ to put them in [increasing order](#). In other words, rearrange the elements of the array so that $v[1] \leq \dots \leq v[n]$.

The other chapter analysed some simple algorithms for the problem. The present chapter examines Heapsort, an algorithm [discovered by J.W.J. Williams](#) in 1964. Unlike the simple algorithms, Heapsort is [linearithmic](#), even in the worst case.

We shall assume that the indices of the array are $1 \dots n$, rather than the usual $0 \dots n-1$. This convention will make the code a little simpler.

Table of contents:

- [Arrays and binary trees](#)
- [Heap](#)
- [Building a heap](#)
- [The sieve function](#)
- [The Heapsort algorithm](#)
- [Animations of Heapsort](#)
- [Performance of Heapsort](#)


Arrays and binary trees

In order to discuss the Heapsort, we must learn to see the binary tree hidden in any array. The set of indices of any array $v[1 \dots m]$ can be understood as a binary tree in the following way:

- the index 1 is the *root* of the tree;
- the *parent* of any index c is $c/2$ (of course 1 has no parent);
- the *left child* of any index p is $2p$ (this child exists only if $2p \leq m$);
- the *right child* of p is $2p+1$ (this child exists only if $2p+1 \leq m$).

1	2	3	4	5	6	7	8	9	10	11	12	13	14
999	999	999	999	999	999	999	999	999	999	999	999	999	999

To make the binary tree stand out, we can draw the array in *layers*, so that every child sits in the layer immediately below that of its parent. The figure below is such a drawing of the array $v[1..56]$. (The numbers in the boxes are the indices i rather than the values $v[i]$.) Observe that each layer, except perhaps the last, has twice as many elements as the previous one. It follows that the number of layers in an array $v[1..m]$ is exactly $1 + \lg(m)$, where $\lg(m)$ is the floor of $\log m$.



Exercises 1

1. Let m be a number of the form $2^k - 1$. Show that more than half of the elements of any array $v[1..m]$ is in the last layer.

Heap

The mechanism behind the Heapsort algorithm is a data structure, known as [heap](#), that sees the array as a binary tree. There are two flavors of the structure: max-heap and min-heap; we shall consider here only the first flavor and omit the “max-” prefix.

A heap, then, is an array in which the value of each parent is greater than or equal to the value of each of its two children. More precisely, an array $v[1..m]$ is a *heap* if

$$v[c/2] \geq v[c]$$

for $c = 2, \dots, m$. Here, as in the rest of this chapter, we shall agree that expressions figuring as indices of arrays are always computed in [integer arithmetic](#). Hence, the value of the expression $c/2$ is $\lfloor c/2 \rfloor$, i.e., the [floor](#) of $c/2$.

1	2	3	4	5	6	7	8	9	10	11	12	13	14
999	888	777	555	666	777	555	222	333	444	111	333	666	333

Occasionally, we consider certain “defective” heaps: we say that an array $v[1..m]$ is a heap *except perhaps for index* k if the inequality $v[c/2] \geq v[c]$ holds for every c distinct from k .

Exercises 2

1. Is the array 161 41 101 141 71 91 31 21 81 17 16 a heap?
2. Show that every [decreasing](#) array is a heap. Show that the converse is not true.
3. ★ Write a function to decide whether an array $v[1..m]$ is a heap.
4. ★ Show that $v[1..m]$ is a heap if and only if each index p has the following properties: (a) $v[p] \geq v[2p]$ provided $2p \leq m$ and (b) $v[p] \geq v[2p+1]$ provided $2p+1 \leq m$.
5. Suppose that $v[1..m]$ is a heap and p is an index smaller than $m/2$. Is it true that $v[2p] \geq v[2p+1]$? Is it true that $v[2p] \leq v[2p+1]$?
6. Suppose that $v[1..m]$ is a heap and let $i < j$ be two indices such that $v[i] < v[j]$. If $v[i]$ and $v[j]$ are interchanged, will $v[1..m]$ still be a heap? Repeat the exercise under the hypothesis $v[i] > v[j]$.
7. ★ Suppose that $v[1..m]$ is a heap and the elements of the array are pairwise distinct. Of course the largest element of the array is $v[1]$. Where can the second largest element be? Where can the third largest element be? Is it true that the third largest element is a child of the second largest element?

Building a heap

It is easy to rearrange the elements of an array $v[1..m]$ of integers so that it becomes a heap. Just repeat the following process: while the value of a child is larger than the value of its parent, swap the values of parent and child and move one step up, towards the root. More precisely, while $v[c/2] < v[c]$, do $\text{swap}(v[c/2], v[c])$ and then $c = c/2$. The swap operation is

```
#define swap (A, B) {int t = A; A = B; B = t;}
```

Here is the complete code:

```
// Rearranges an array v[1..m] so that
// it becomes a heap.

static void
buildheap (int m, int v[])
{
    for (int k = 1; k < m; ++k) {
        // v[1..k] is a heap
        int c = k+1;
        while (c > 1 && v[c/2] < v[c]) { // 5
            swap (v[c/2], v[c]); // 6
            c /= 2; // 7
        }
    }
}
```

(The keyword [static](#) indicates that `buildheap` is an auxiliary function that cannot be called directly by the user of Heapsort.)

At the beginning of every iteration controlled by the “for”, the array $v[1..k]$ is a heap. In the course of the iteration, $v[k+1]$ moves up the heap (towards the root) until it finds its correct place and is thus incorporated into the heap.

In each repetition of the pair of lines 6-7, the index c jumps from one [layer](#) of the array to the previous layer. Hence, this pair of lines can be repeated at most $\lg(k)$ times for each fixed k . As a consequence, the total number of comparisons (line 5 of the code) between elements of the array is at most

$$m \lg(m) .$$

(As we shall see [further down](#), it is possible to do the job with just m comparisons.)

Exercises 3

1. IMPORTANT. Criticize the following idea: to transform an array into a heap, just rearrange it in decreasing order (using [Mergesort](#) or [Quicksort](#), for example).
2. Prove that function `buildheap` is correct. Begin by establishing the [invariants](#) of the iterative process that span the lines 5 to 7.
3. Discuss the following version of the `buildheap` function:

```
for (int k = 1; k < m; ++k) {
    int c = k+1, x = v[k+1];
    while (c > 1 && v[c/2] < x) {
        v[c] = v[c/2];
        c /= 2; }
    v[c] = x; }
```

The sieve function

The core of many algorithms that manipulate heaps is a function that, unlike [buildheap](#), moves *down* the heap, away from the root. This function, which we call `sieve`, receives an arbitrary array $v[1..m]$ and

moves $v[1]$ down to its correct position,

jumping from one layer to the next. How is this done? If $v[1] \geq v[2]$ and $v[1] \geq v[3]$, nothing needs to be done. If $v[1] < v[2]$ and $v[2] \geq v[3]$, just swap $v[1]$ with $v[2]$ and move $v[2]$ down to its correct position. In the other two cases, do something similar. (See a [draft](#) of the algorithm in pseudocode.) In the following example, each line of the figure shows the state of the array at the beginning of an iteration:

85	99	98	97	96	95	94	93	92	91	90	89	97	86
99	85	98	97	96	95	94	93	92	91	90	89	97	86
99	97	98	85	96	95	94	93	92	91	90	89	97	86
99	97	98	93	96	95	94	85	92	91	90	89	97	86

We can now write the code of the function. Each iteration begins with an index p and chooses a child c of p which has the largest value:

```
static void
sieve (int m, int v[]) {
    int c = 2;
    while (c <= m) {
        if (c < m && v[c] < v[c+1]) ++c;
        // c is a most valuable child of c/2
        if (v[c/2] >= v[c]) break;
        swap (v[c/2], v[c]);
```

```

        c *= 2;
    }
}

```

The function will be applied to arrays that are heaps [except perhaps](#) for one or two indices. The function can, therefore, be documented as follows:

```

// Receives an array v[1..m] that is a heap
// except perhaps for indices 2 and 3 and
// rearranges the array so that it becomes
// a heap.

```

The following version is a little better, because it moves fewer elements of the array from one place to another (and does fewer divisions of c by 2):

```

static void
sieve (int m, int v[])
{
    int p = 1, c = 2, x = v[1];
    while (c <= m) {
        if (c < m && v[c] < v[c+1]) ++c;
        if (x >= v[c]) break;
        v[p] = v[c];
        p = c, c = 2*p;
    }
    v[p] = x;
}

```

Performance. The sieve function is very fast. It does at most $\lg(m)$ iterations, since the array has $1 + \lg(m)$ [layers](#). Each iteration involves two comparisons between elements of the array and therefore the total number of comparisons is at most

$$2 \lg(m) .$$

The *time* spent is proportional to the number de comparisons and therefore proportional to $\log m$ in the worst case.

Exercises 4

1. Why is the following implementation of the sieve function incorrect?

```

int p = 1, c = 2;
while (c <= m) {
    if (v[p] < v[c]) {
        swap (v[p], v[c]);
        p = c;
        c = 2*p; }
    else {
        if (c < m && v[p] < v[c+1]) {
            swap (v[p], v[c+1]);
            p = c+1;
            c = 2*p; }
        else break; } }

```

2. Is the following alternative code of the sieve function correct?

```

for (int c = 2; c <= m; c *= 2) {
    if (c < m && v[c] < v[c+1]) ++c;
    int p = c/2;
}

```

```

    if (v[p] >= v[c]) break;
    swap (v[p], v[c]); }

```

3. Discuss the following variant of the code of `sieve`:

```

int x = v[1], c = 2;
while (c <= m) {
    if (c < m && v[c] < v[c+1]) ++c;
    if (x >= v[c]) break;
    v[c/2] = v[c];
    c *= 2; }
v[c/2] = x;

```

4. RECURSIVE. Write a recursive version of the function `sieve`.

5. ★ GENERALIZED SIEVE. Suppose that an array $v[1..m]$ is a heap except perhaps for indices $2p$ and $2p+1$. Write a function `sieve2` that will receive $v[1..m]$ and p and transform the array into a heap.

6. ★ BUILDING A HEAP FASTER. Show that the following function has the same effect as [buildheap](#), that is, it transforms array $v[1..m]$ into a heap:

```

void buildheap2 (int m, int v[]) {
    for (int p = m/2; p >= 1; --p)
        sieve2 (p, m, v);
}

```

Show that `buildheap2` does at most $(m \lg(m))/2$ comparisons between elements of the array. Refine your analysis to show that the function actually does at most m comparisons.

7. Does the following code fragment transform array $v[1..m]$ into a heap?

```

for (int p = 1; p <= m/2; ++p)
    sieve2 (p, m, v);

```

8. PRIORITY QUEUE. A [priority_queue](#) is a set of n objects (numbers, for example) subject to two operations: (1) deletion of a largest element and (2) insertion of a new object. If the set is kept in a heap, these two operations can be made very fast. Implement the two operations so that each takes time proportional to $\log n$ in the worst case.

The Heapsort algorithm

We can now put together all the pieces discussed above and write an algorithm that will rearrange an array $v[1..n]$ in increasing order. The algorithm has two phases: the first transforms the array into a heap and the second pulls elements from the heap in decreasing order. (See a [draft](#) of the algorithm.)

```

// Rearranges the elements of array v[1..n]
// in increasing order.

void
heapsort (int n, int v[])
{
    buildheap (n, v);
    for (int m = n; m >= 2; --m) {
        swap (v[1], v[m]);
        sieve (m-1, v);
    }
}

```

At the beginning of each iteration of the “for”, the following [invariant](#) properties hold:

- $v[1..m]$ is a heap,
- $v[1..m] \leq v[m+1..n]$,
- $v[m+1..n]$ is in increasing order, and

- $v[1..n]$ is a permutation of the original array.


The expression “ $v[1..m] \leq v[m+1..n]$ ” is a shorthand for “each element of $v[1..m]$ is smaller than or equal to every element of $v[m+1..n]$ ”.

1	heap							m	increasing				n
888	777	666	555	444	333	222	111	000	999	999	999	999	999
small elements								large elements					

It follows that $v[1..n]$ will be in increasing order when m becomes equal to 1. This shows that the algorithm is correct.

Animations of Heapsort

The animation at right (copied from [Simon Waldherr / Golang Sorting Visualization](#)) shows Heapsort running on an array $v[0..79]$ of positive numbers. Each element $v[i]$ of the array is represented by a point $(i, v[i])$. (For some reason, the animation does not execute the last two iterations.)



Here is a sample of other animations:

- [Heapsort](#), Wikipedia entry.
- [Heapsort animation](#), by David Galles (University of San Francisco).
- Animation of [15 sorting algorithms](#), by Timo Bingmann, on YouTube.
- [Sorting Algorithms Animations](#) on Toptal.

Exercises 5

1. Use the `heapsort` function to sort the array 16 15 14 13 12 11 10 9 8 7 6 5 4 . Show the state of the array at the beginning of each iteration.
2. Describe the state of the array at the beginning of a generic iteration of the Heapsort algorithm?
3. CORRECTNESS CHECK. Write a program to test, experimentally, the correctness of your implementation of the Heapsort algorithm. (See the [analogous exercise for Insertionsort](#).)
4. Does the `heapsort` function produce a [stable](#) rearrangement of the array? In other words, does it preserve the relative order of elements that have the same value?
5. Write a function with prototype `heap_sort (int *v, int n)` that will rearrange an array `v[0..n-1]` (notice the indices) in increasing order. (This can be done by an appropriate call to the [heapsort](#) function.)
6. MIN-HEAP. Write a function that will rearrange an array `v[1..n]` in *decreasing* order. (Hint: Adapt the definition of `heap` and rewrite the `sieve` function.)
7. Emulate a heap by means of a set of cells interconnected by [pointers](#). Each cell will have four fields: a `value` and three pointers, one pointing the parent of the cell, another pointing the right child, and another pointing the left child. Write an appropriate version of the `sieve` function. (See the [Binary trees](#) chapter.)

Performance of Heapsort

How many comparisons between elements of the array does the [heapsort](#) function execute? The auxiliary function `buildheap` does at most [\$n \lg\(n\)\$ comparisons](#). Next, the `sieve` function is called approximately n times and each of these calls does at most [\$2 \lg\(n\)\$ comparisons](#). Hence, the total number of comparisons is at most

$$3 n \lg(n).$$

The *time* consumed by heapsort is proportional to the number of comparisons between elements of the array, and therefore proportional to $n \log n$ in the worst case. (But the proportionality factor is larger than that of [Mergesort](#) and [Quicksort](#).)

Exercises 6

1. PERFORMANCE TEST. Write a program to time your implementation of the Heapsort algorithm. (See the [analogous exercise for Mergesort](#).)

See chapter 14 of the [Programming Pearls](#) by Jon Bentley.

Updated 2019-02-02

<https://www.ime.usp.br/~pf/algorithms/>

© [Paulo Feofiloff](#).

[CS-IME-USP](#)